

The Quantum Random Oracle Toolbox

Hans Heum 

NTNU - Norwegian University of Science and Technology, Trondheim, Norway.
`hans.heum@ntnu.no`

Abstract. In which we reveal the tricks of the trade.

Keywords: Provable Security · Quantum Random Oracles

Table of Contents

References	3
1 Introduction.....	3
2 Preliminaries	3
2.1 A Brief Quantum Primer	3
3 Defining Quantum Random Oracles	5
4 Reprogramming QROs: The O2H Lemma and Friends	6
4.1 Recording ROs (or, Random Oracles are Good PRFs).....	6

1 Introduction

Hans: Idea: Take the PKE from “Random Oracles Are Practical” [BR93] as a motivating example running through the paper, using stronger and stronger techniques to show stronger and stronger security results.

There are three main challenges for proofs in the quantum random oracle model, relating to the application of standard techniques in the classical random oracle model:

1. Recording: or, how to extract QRO queries.
2. Reprogramming: or, how to alter specific outputs of the QRO.
3. Rewinding: or, how to derandomize quantum oracle algorithms.

2 Preliminaries

Basic notation. We take the natural numbers $\mathbb{N}$ to include 0. For $n \in \mathbb{N}$, we define the set $[n] = \{0, \dots, n-1\}$. We will employ the convention that $N = 2^n$; thus, $NM = 2^{n+m}$.

Functions are defined over bitstrings unless otherwise noted, and the set of all functions from the set of m -length bitstrings $\{0, 1\}^m$ to the set of n -length $\{0, 1\}^n$ bitstrings is denoted $\mathcal{F}_{m \rightarrow n}$.

2.1 A Brief Quantum Primer

Quantum theory arose as a branch of physics, and so much of the standard notational conventions are hand-wavy, and often inconsistently implemented across the literature. For the sake of clarity and self-containedness, we specify all our conventions below, providing a brief summary of the building blocks that make up quantum computing theory along the way; for more details, see the celebrated introductory text by Nielsen and Chuang [1].

Provide explanations of:

1. Bare-bones background on quantum states and Dirac notation; states vs. labels.
2. $N = 2^n$; the set $[n]$; n -length registers are described by vectors in an N -sized Hilbert space.
3. Greek letters: When and why I use them.
4. Precise notational conventions regarding registers.
5. Precise conventions regarding fonts.
6. Notation for unitary evolutions: $\xrightarrow{U}$.
7. Stick to pure states, and only mention the generalization to density matrices.
8. Basic background on game hopping and advantages.
9. Maybe write something about asymptotic vs. concrete security.

Dirac notation. In general, quantum states are complex-valued vectors in a Hilbert space (a vector space defined over functions), which may or may not be finite-dimensional.¹ Quantum states are represented using the Dirac notation, in which a quantum state carrying the *label* ψ is written $|\psi\rangle$, its conjugate transpose is written $\langle\psi| := |\psi\rangle^\dagger := (|\psi\rangle^*)^T$, and its inner product with itself is written $\langle\psi|\psi\rangle := ||\psi||^2$. Quantum states are unit-norm, and so we have that $\langle\psi|\psi\rangle = 1$. For two not necessarily equal quantum states $|\psi\rangle$ and $|\varphi\rangle$, their inner product $\langle\psi|\varphi\rangle$ is a real number between zero and one, where “zero” means the states are orthogonal, and “one” means the states are identical (modulo an unobservable global phase factor $e^{i\theta}$).

Qubit basis states. We are interested in finite-dimensional quantum states, implemented on a collection of n qubits. A *qubit* is a quantum state existing in the two-dimensional Hilbert space characterized by the basis states $|0\rangle$ and $|1\rangle$. In physical implementations, these states may for example be mapped to the spin of an electron (“up”/“down”), or the two lowest energy states of an atom, or in general any two orthogonal quantum states, where two quantum states are orthogonal if and only if there exists a measurement procedure that perfectly distinguishes them.

¹ For example, consider the state $|p(x)\rangle$ describing the position $x \in [0, 1]$ of particle p ; this will live in an infinite-dimensional Hilbert space characterized by the basis vectors $\{|p(x)\rangle\}_{x \in [0, 1]}$.

It is worth emphasizing that the “0” and “1” here are *labels*, and do not say anything about the physical system in question other than that it contains (at least) two orthogonal quantum states, which are being identified with the two basis states $|0\rangle$ and $|1\rangle$.

Since two qubit basis states $|b_1\rangle$ and $|b_2\rangle$, for $b_1, b_2 \in \{0, 1\}$, the two states must be either identical (up to the aforementioned phase factor) or orthogonal, we have that

$$\langle b_1 | b_2 \rangle = \delta_{b_1 b_2} ,$$

where δ_{ij} is the Kronecker delta, equal to 1 if $i = j$, and 0 otherwise.

Combined quantum states. In general, two quantum systems (i.e. states describing distinct physical systems) may be combined via the tensor product:

$$|\psi\rangle \otimes |\varphi\rangle := |\psi\varphi\rangle .$$

If $|\psi\rangle$ describes an N -dimensional quantum system, and $|\varphi\rangle$ describes an M -dimensional quantum system, then the combined state $|\psi\varphi\rangle$ describing the combined system will be NM -dimensional. (Note how the dimensions are *multiplied* rather than added, as one might naïvely expect.)

Applying this to qubits, given a collection of n qubits $|b_i\rangle$, each initialized to one of the two basis states (so that $b_i \in \{0, 1\}$), we may construct the combined state

$$|b_0 b_1 \dots b_{n-1}\rangle := |b_0\rangle \otimes |b_1\rangle \otimes \dots \otimes |b_{n-1}\rangle .$$

Since each qubit is a two-dimensional quantum state, we see that the resulting combined state will have dimension $N := 2^n$. The resulting Hilbert space will be characterized by the computational basis states²

$$\{|i\rangle\}_{i \in [N]} = \{|b_0 b_1 \dots b_{n-1}\rangle\}_{\forall b_0, \dots, b_{n-1} \in \{0, 1\}} .$$

This exponentiality in the dimension of quantum states as the number of qubits increases sits at the core of the quantum advantage over classical computers. As an illustrative example, the Hilbert space corresponding to a three-qubit quantum system is characterized by the eight computational basis states:

$$\begin{aligned} \{|i\rangle\}_{i \in [2^3]} &= \{|000\rangle, |001\rangle, |010\rangle, |011\rangle, |100\rangle, |101\rangle, |110\rangle, |111\rangle\} \\ &= \left\{ \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \end{bmatrix} \right\} . \end{aligned}$$

As in the single-qubit case, we see that we have orthonormality over the computational basis states:

$$\langle x | y \rangle = \delta_{xy} .$$

Qubit registers. Given a collection of n qubits, we will often find it useful to divide it into *registers* of specified sizes, for example to separate distinct elements in an input. We will use subscripts to refer to specific registers, unless the register we are referring to is clear from context.

As an example, let’s say we want the first n_1 qubits store the variable x —call this “Register 1”—and that the remaining n_2 qubits—“Register 2”—should store the variable y . Then the following are all equivalent ways of denoting the combined state:

$$|x||y\rangle = |x\rangle_1 \otimes |y\rangle_2 := |x\rangle_1 |y\rangle_2 := |x, y\rangle ,$$

where $\|$ denotes string concatenation.

² As with any vector space, Hilbert spaces permit an infinitude of basis choices; the “computational” basis is simply the name of the choice wherein basis states are labelled by classical bits (or, in general, bitstrings).

Greek and Roman letters. We will use the letters $x, y, z, \dots$ to represent classical binary strings. Thus, a state $|x\rangle$ will always be understood to be a computational basis state. Whenever we are talking specifically about single-qubit computational basis states, we will use b to represent a single bit, and write $|b\rangle$.

General quantum states are traditionally labelled by letters taken from the middle of the greek alphabet, such as ψ, φ, ξ , etc. We will follow this tradition, or otherwise use a descriptive label relating to the specific context in which the state is being used.

Quantum *amplitudes*, meanwhile, are written using letters taken from the start of the greek alphabet ($\alpha, \beta, \dots$).

General qubit states. A general single-qubit state may be written in terms of the computational basis states as follows:

$$|\psi\rangle = \alpha_0 |0\rangle + \alpha_1 |1\rangle ,$$

where the amplitudes α_i are complex numbers satisfying $|\alpha_0|^2 + |\alpha_1|^2 = 1$, (where the absolute value of a complex number α is defined as $|\alpha| = \sqrt{\alpha^* \alpha}$).

Whenever a qubit is in a state described by amplitudes α_0 and α_1 that are both non-zero, we say that the qubit is in a *superposition* over the values “0” and “1”. If we then *observe* the qubit, the probability that the observed outcome X is either 0 or 1 is then given by $\Pr[X = b] = |\alpha_b|^2 = |\langle\psi|b\rangle|^2$, and after observation, the superposition will have “collapsed” to the basis state corresponding to the observed value, $|b\rangle$.³

Likewise, a general, n -qubit quantum state may be written in terms of the computational basis states:

$$|\psi\rangle = \sum_{\forall i \in [2^n]} \alpha_i |i\rangle = \sum_{i=0}^{N-1} \alpha_i |i\rangle = \alpha_0 |0\dots 00\rangle + \alpha_1 |0\dots 01\rangle + \dots + \alpha_{N-1} |1\dots 11\rangle ,$$

where the *amplitudes* α_i again satisfy $\sum_{\forall i \in [2^n]} |\alpha_i|^2 = 1$.

When expanding a general quantum state in terms of the computational basis states in the manner above, we will typically exclude terms for which $\alpha_i = 0$, writing instead

$$|\psi\rangle = \sum_i \alpha_i |x_i\rangle := \sum_{i \in \mathcal{I}} \alpha_i |x_i\rangle ,$$

where $\mathcal{I}$ indexes the set $\{|x\rangle \mid x \in [N] \wedge \langle\psi|x\rangle \neq 0\}$.

Hans: Give inner product as Euclidean distance over amplitudes.

Unitary operations.

Hans: Next.

3 Defining Quantum Random Oracles

A classical oracle $\mathcal{O}_f$ is a black-box procedure that, given input x , returns $f(x)$ in such a way that no information about f other than that which can be inferred from its output is provided. A classical

³ The discernible effect of observations on quantum states has led to a century of heated philosophical debate over how quantum theory should be interpreted: The term “collapse” used here is typically associated with the conservative “Copenhagen interpretation”, championed by Niels Bohr, and commonly characterized as taking a “shut-up-and-calculate” attitude towards the problem. Another popular interpretation is the many-worlds interpretation due to Hugh Everett, and favoured by father of quantum computing theory David Deutsch, among others. Instead of the quantum state “collapsing” to become “classical”, this interpretation would say that both terms are ill-defined, and that the universe should instead be understood to be nothing but one big superposition state over all worlds that might have arisen from the initial conditions of the Big Bang; and that what actually happened when you made your observation, was that the quantum state describing the subsystem you were studying and the quantum state describing the subsystem containing *you* became correlated via the process known as *entanglement*. (In case you’re wondering, I place my own bet on the *relational* interpretation due to Carlo Rovelli, which may be described as “many-worlds without the multiverse”.)

random oracle, then, is simply a classical oracle where the function f was chosen uniformly at random from the space of all functions from $\{0, 1\}^m$ to $\{0, 1\}^m$.⁴

However, the above operation is not unitary, as required for evolutions of quantum states (which is what we mean by “quantum algorithms”). Therefore, any quantum algorithm must instead be given access to *quantum-accessible* oracles, which on input the computational basis state $|x, y\rangle$, returns the unitary evolution of the state $|x, y \oplus f(x)\rangle$. In general (letting n and m be the length of the x and y registers, respectively), on input a quantum state $|\psi\rangle = \sum_{i \in [2^{n+m}]} \alpha_i |x, y\rangle$, we get:

$$\sum_{i \in [2^{n+m}]} \alpha_i |x_i, y_i\rangle \xrightarrow{O_f} \sum_{i \in [2^{n+m}]} \alpha_i |x_i, y_i \oplus f(x_i)\rangle .$$

To see that this operation is unitary, it is enough to observe that it is *reversible*: Inputting the same state a second time yields the original state:

$$\sum_{i \in [2^{n+m}]} \alpha_i |x_i, y_i \oplus f(x_i)\rangle \xrightarrow{O_f} \sum_{i \in [2^{n+m}]} \alpha_i |x_i, y_i \oplus f(x_i) \oplus f(x_i)\rangle = \sum_{i \in [2^{n+m}]} \alpha_i |x_i, y_i\rangle .$$

Sometimes it we will want to provide a quantum algorithm with access to a classical oracle.

Hans: Next.

4 Reprogramming QROs: The O2H Lemma and Friends

4.1 Recording ROs (or, Random Oracles are Good PRFs)

Hans: Copy-pasted from first intro draft. *Maybe* do the following later when we introduce the compressed random oracle technique? But definitely do not *start* here. (Game figs. in comments.) (Is O2H applicable to this though? Maybe we could provide proofs using both techniques?)

We start with the simplest ROM proof imaginable: showing that random oracles are good pseudo-random functions. Specifically, we will show that $\text{PRF}_k(x) = H(k||x)$ is indistinguishable from a random function, provided the key k is hard to guess.

We will do this by showing that if an adversary $\mathbb{A}$ can distinguish the output of PRF_k from random, then a reduction $\mathbb{B}$ can recover the key k . Crucially, the reduction does this by *recording* the queries of the adversary $\mathbb{A}$ to the random oracle H .

⁴ Typically, random oracles follow hash functions in being defined from the space of all finite-length bit strings $\{0, 1\}^*$ to the space of n -length bitstrings $\{0, 1\}^n$; technically, this simply means that f is a collection of random functions: one per input length. Since quantum random oracles must operate by unitary means, we will explicitly restrict the domain of the random function to $\{0, 1\}^m$.

References

- BR93. Mihir Bellare and Phillip Rogaway. Random oracles are practical: A paradigm for designing efficient protocols. In Dorothy E. Denning, Raymond Pyle, Ravi Ganesan, Ravi S. Sandhu, and Victoria Ashby, editors, *ACM CCS '93*, pages 62–73. ACM Press, November 1993.